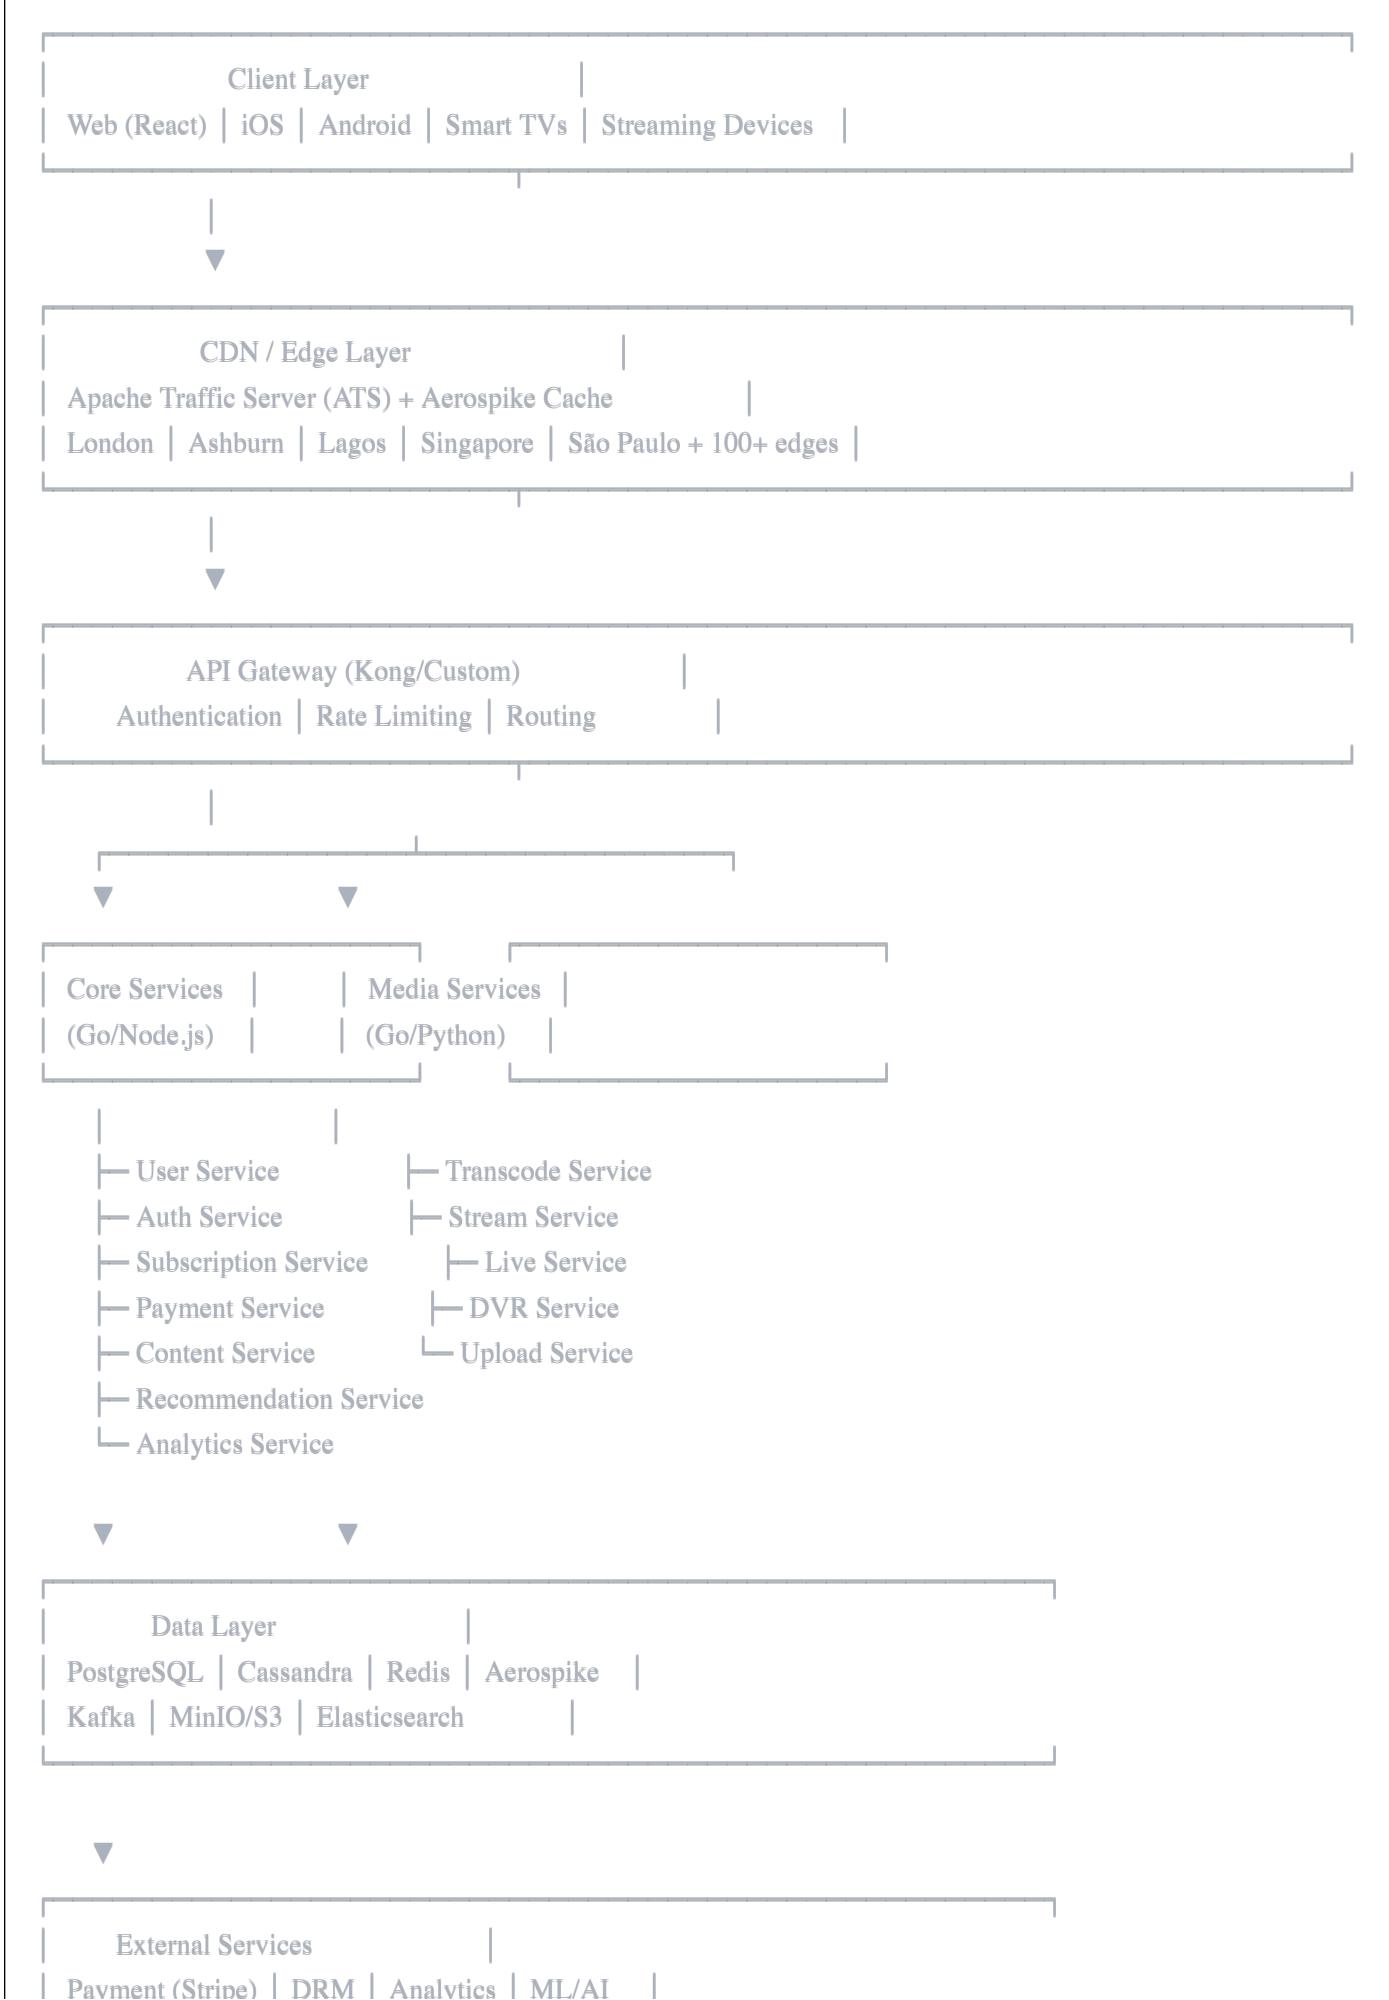


System Architecture: Global Streaming Platform

1. High-Level Architecture



2. Microservices Architecture

2.1 Core Services

User Service

- **Responsibility:** User profiles, preferences, watch history
- **Tech:** Go + PostgreSQL + Redis
- **API:** gRPC internal, REST external
- **Scale:** 10K RPS per instance

Auth Service

- **Responsibility:** Authentication, JWT tokens, sessions
- **Tech:** Go + Redis + PostgreSQL
- **Features:** OAuth 2.0, JWT, refresh tokens, MFA
- **Scale:** 50K RPS per instance

Subscription Service

- **Responsibility:** Plans, billing cycles, entitlements
- **Tech:** Go + PostgreSQL + Kafka
- **Integration:** Stripe webhooks, payment processing
- **Scale:** 5K RPS per instance

Content Service

- **Responsibility:** Metadata, catalogs, search
- **Tech:** Go + PostgreSQL + Elasticsearch
- **Features:** Multi-language metadata, rights management
- **Scale:** 100K RPS per instance (read-heavy)

Recommendation Service

- **Responsibility:** Personalized recommendations
- **Tech:** Python + TensorFlow Serving + Redis
- **Models:** Collaborative filtering, deep learning
- **Scale:** 20K RPS, <100ms P95 latency

2.2 Media Services

Transcode Service

- **Responsibility:** VOD transcoding pipelines
- **Tech:** Go orchestration + GStreamer workers
- **Profiles:** 6-8 ABR profiles (240p-4K)
- **Scale:** 1000+ concurrent jobs

Stream Service

- **Responsibility:** Manifest generation, packaging
- **Tech:** Go + Nginx/ATS
- **Protocols:** HLS, DASH, LL-HLS
- **Scale:** 1M+ concurrent streams

Live Service

- **Responsibility:** Live ingest, transcoding, distribution
- **Tech:** OvenMediaEngine + Go orchestration
- **Features:** WebRTC, LL-HLS, DVR buffer
- **Scale:** 10K+ concurrent live streams

DVR Service

- **Responsibility:** Cloud DVR recording, playback
- **Tech:** Go + GStreamer + MinIO
- **Storage:** Time-shifted content (7-30 days)
- **Scale:** 100K concurrent recordings

2.3 Analytics & ML Services

Analytics Service

- **Responsibility:** Event ingestion, real-time metrics
- **Tech:** Go + Kafka + Cassandra + Druid
- **Events:** Play, pause, seek, buffering, errors
- **Scale:** 1M events/sec

ML Service

- **Responsibility:** Model training, inference
- **Tech:** Python + PyTorch + Kubeflow
- **Models:**
 - Recommendation (collaborative + deep learning)
 - CDN selection (gradient boosting)
 - Churn prediction (XGBoost)
 - Content understanding (computer vision)

3. Data Architecture

3.1 Databases

PostgreSQL (Transactional)

- **Usage:** User accounts, subscriptions, content metadata
- **Setup:** Primary-replica with read replicas per region
- **Backup:** Daily full, hourly incremental (30-day retention)
- **HA:** Patroni + pgBouncer connection pooling

Cassandra (Time-Series & High Write)

- **Usage:** Analytics events, playback history, logs
- **Setup:** 3 nodes per datacenter, RF=3
- **Compaction:** Time-window compaction strategy
- **TTL:** 90 days for raw events, 2 years for aggregates

Redis (Cache & Session)

- **Usage:** Session storage, hot data cache, rate limiting
- **Setup:** Redis Cluster (6 shards, 3 replicas each)
- **Eviction:** LRU for cache, no eviction for sessions
- **Persistence:** AOF for sessions, snapshots for cache

Aerospike (CDN Cache Metadata)

- **Usage:** Content hotness scores, TTL hints, edge cache coordination
- **Setup:** 5 nodes per Tier 1 PoP (rack-aware)
- **Replication:** Cross-region replication for hotness scores
- **Performance:** <1ms P99 read latency

3.2 Message Queue

Apache Kafka

- **Topics:**
 - `(user.events)` (clicks, views)
 - `(playback.events)` (play, pause, seek)
 - `(billing.events)` (subscriptions, payments)
 - `(content.events)` (uploads, transcode jobs)
- **Retention:** 7 days (events), 30 days (billing)
- **Partitions:** 50-100 per topic (based on throughput)
- **Consumers:** Analytics, ML training, real-time dashboards

3.3 Object Storage

MinIO (Warm/Cold Tier)

- **Usage:** Transcoded video segments, images, backups
- **Setup:** Distributed mode, erasure coding (EC:4+2)
- **Lifecycle:** Auto-tier to cold storage after 60 days
- **CDN:** Origin for ATS cache pulls

Cloudflare R2 (Global Hot Tier)

- **Usage:** Most-watched content globally
- **Setup:** R2 + Cache Reserve for ultra-low latency
- **Sync:** Selective replication based on Aerospike hotness scores

4. CDN Architecture

4.1 Apache Traffic Control (ATC) Setup

Traffic Ops (Control Plane)

- **Function:** CDN configuration, delivery service management
- **API:** RESTful API for automation
- **Database:** PostgreSQL for CDN config
- **UI:** Web console for operators

Traffic Router (DNS/HTTP Routing)

- **Function:** Geo-based user → optimal edge routing
- **Method:** GeoDNS + HTTP 302 redirects
- **Metrics:** Real-time latency, capacity monitoring
- **Failover:** Automatic unhealthy edge exclusion

Traffic Monitor

- **Function:** Edge server health checks, metrics collection
- **Interval:** 10s health checks, 60s stats aggregation
- **Data:** CPU, disk, network, cache hit ratio per edge
- **Alerting:** PagerDuty integration for failures

4.2 Apache Traffic Server (ATS) Configuration

L2 Parent Servers (Tier 1 PoPs)

- **Hardware:** 128GB RAM, 20TB NVMe SSD, 40Gbps NIC
- **Cache Size:** 15TB hot cache per server
- **Config:**
 - Strict parenting: `(go_direct=false)`
 - Origin: MinIO, R2
 - TTL: Manifest 15s, segments 3600s (from Aerospike)

L1 Edge Servers (100+ Locations)

- **Hardware:** 64GB RAM, 4TB NVMe SSD, 10Gbps NIC
- **Cache Size:** 3TB hot cache per server
- **Config:**
 - Parent: Nearest L2
 - Cachekey plugin: Remove `?token=` query params
 - Header rewrite: Add `(X-Region)`, remove `(Pragma)`

4.3 Aerospike Integration

Hotness Tracking

Key: content_id (video/channel ID)

Record: {

```
    hotness_score: float (0-1, computed from recent views),  
    ttl_hint: int (seconds, for ATS cache),  
    last_accessed: timestamp,  
    region_scores: map<region, float>  
}
```

Replication Strategy

- **S1 (Tier 0):** Top 1% hottest content → replicate globally
- **S2 (Tier 1):** Top 10% → replicate to relevant regions
- **S3 (Long Tail):** Pull from origin on-demand, cache locally

TTL Management (Lua Plugin in ATS)

lua

```
function do_remap()  
    local content_id = extract_content_id(ts.client_request.get_uri())  
    local aerospike_record = query_aerospike(content_id)  
  
    if aerospike_record then  
        local ttl = aerospike_record.ttl_hint  
        ts.client_response.header['Cache-Control'] = 'max-age=' .. ttl  
    else  
        -- Default TTL  
        ts.client_response.header['Cache-Control'] = 'max-age=3600'  
    end  
  
    return 0  
end
```

5. Streaming Architecture

5.1 VOD Workflow

1. Content Upload → MinIO
2. Transcode Job → Kafka → Transcode Service
3. GStreamer: Generate ABR ladder (6 profiles)
4. Package: HLS/DASH manifests + segments → MinIO
5. Manifest URL → Content Service → User request
6. Playback:
 - User → DNS → Traffic Router
 - Traffic Router → Optimal ATS edge
 - ATS edge → (cache hit?) → User
 - ATS edge → (cache miss?) → L2 parent → MinIO → cache → User

5.2 Live Streaming Workflow

1. Ingest Source (RTMP/SRT) → OvenMediaEngine cluster
2. OME: Transcode to ABR profiles (real-time)
3. OME: Generate LL-HLS (2s segments, 0.5s parts)
4. OME: Publish to ATS edge (origin-push or edge-pull)
5. Playback:
 - User → ATS edge (cached live segments)
 - <3s glass-to-glass latency

5.3 FAST Channel Architecture

1. Scheduler Service: Pre-planned content grid (24hrs+)
2. Playout Service:
 - Read schedule → fetch VOD segments from cache
 - Generate continuous HLS stream (server-side stitching)
 - Insert ads (SSAI) based on ad decision service
3. Output: HLS/DASH endpoint per channel
4. Distribution: Via ATS CDN (same as VOD)

6. Security Architecture

6.1 DRM Flow

1. User requests protected content
2. Auth Service: Verify entitlement (subscription/purchase)
3. License Server:
 - Generate license (Widevine/PlayReady/FairPlay)
 - Include playback restrictions (HDR, 4K based on tier)
4. Client: DRM decryption → playback
5. Content: AES-128 or AES-256 encrypted segments

6.2 API Security

- **Authentication:** JWT (access token 15min, refresh token 30 days)
- **Authorization:** RBAC (roles: admin, editor, viewer, guest)
- **Rate Limiting:** Token bucket (100 req/min per user, 10K/min per API key)
- **WAF:** Cloudflare/AWS WAF for DDoS, injection attacks
- **Encryption:** TLS 1.3 only, HSTS enabled

7. Scalability & Performance

7.1 Auto-Scaling Strategy

Service	Metric	Scale Up	Scale Down
API Gateway	CPU > 70%	+20% instances	-10% instances
Stream Service	Concurrent streams > 80% capacity	+50% instances	-25% instances
Transcode Workers	Queue depth > 100 jobs	+10 workers	-5 workers
Database Read Replicas	Replication lag > 5s	+1 replica	Manual

7.2 Performance Targets

Metric	Target	Measurement
API Response Time (P95)	<200ms	Prometheus
Video Startup Time (P95)	<2s	Client telemetry
Buffering Ratio	<0.5%	Playback events
CDN Hit Ratio	≥90%	ATS logs → Aerospike
Concurrent Streams	5M+	Stream Service metrics
Transcode Time	1x realtime (VOD)	Transcode Service

7.3 High Availability

- **Multi-AZ Deployment:** All services in 3 AZs per region
- **Database:** Primary + sync replica + 2 async replicas
- **CDN:** N+1 redundancy (if 1 PoP fails, reroute to next nearest)
- **Kafka:** 3 brokers per cluster, RF=3, min ISR=2
- **Load Balancers:** Active-active ALB/NLB with health checks

8. Monitoring & Observability

8.1 Metrics (Prometheus)

Infrastructure:

- CPU, memory, disk, network per service
- Request rate, latency, error rate (RED method)

Business:

- Concurrent streams, new subscribers, churn rate
- Revenue (subscriptions, ads, transactions)

Quality:

- Video startup time, buffering events, bitrate distribution
- CDN hit ratio, origin traffic

8.2 Logs (ELK Stack)

- **Structured Logging:** JSON format with trace IDs
- **Centralized:** Filebeat → Logstash → Elasticsearch
- **Retention:** 30 days (hot), 90 days (warm), 1 year (cold)
- **Alerting:** ElastAlert for error rate spikes

8.3 Tracing (Jaeger/OpenTelemetry)

- **Instrumentation:** All Go services with OpenTelemetry SDK
- **Sampling:** 100% errors, 1% successful requests
- **Storage:** Elasticsearch backend
- **UI:** Jaeger UI for distributed tracing

8.4 Dashboards (Grafana)

- **CDN Ops:** Hit ratio, bandwidth, top content
- **Platform Health:** Service status, error rates, latencies
- **User Engagement:** MAU, watch time, top genres
- **Revenue:** Subscriptions, ads, PPV by region

9. Deployment Architecture

9.1 Kubernetes Setup

- **Clusters:** 1 per Tier 1 PoP (5 total) + 1 management cluster
- **Namespaces:** `(production)`, `(staging)`, `(monitoring)`
- **Ingress:** Nginx Ingress Controller with TLS termination
- **Service Mesh:** Istio for traffic management, observability

9.2 CI/CD Pipeline



9.3 Infrastructure as Code

- **Terraform:** AWS/GCP resources (VPC, databases, load balancers)
- **Ansible:** Server configuration, ATS setup
- **Helm:** Kubernetes app deployments
- **GitOps:** All config in Git, ArgoCD for sync

10. Disaster Recovery

10.1 RTO & RPO Targets

Service	RTO	RPO	Strategy
API Services	15 min	0	Multi-AZ, instant failover
Database	1 hour	5 min	Point-in-time recovery
Content Library	4 hours	24 hours	S3 cross-region replication
CDN	5 min	0	Multi-PoP, auto-rerouting

10.2 Backup Strategy

- **Databases:** Daily full + 15-min incremental (30-day retention)
 - **Content:** 3-2-1 rule (3 copies, 2 media types, 1 offsite)
 - **Config:** Git repositories (immutable history)
-

Document Version: 1.0

Last Updated: 2025-01-08

Architect: Platform Engineering Team